



# Multi-Agent Systems Project Presentation

Professors: Andrea Omicini, Roberta Calegari

Student: Andrea Perna

University of Bologna



# Project Motivations



# Project Goals




## Multi-Agent System

Design a modular Multi-Agent System (MAS) for Search and Rescue missions (SAR) in Python;

Engineer a clear separation between agents and environment;

Create a distributed network of robotic agents modelled with a strongly connected graph;



## Distributed Optimization

Implement distributed aggregative optimization algorithm to provide distributed coordination;

Agents aim to fulfill both local and global targets, by reaching consensus on global information;

Ensure efficient navigation with obstacle avoidance;



## Beliefs-Desires-Intentions

Leverage the BDI framework to govern robots' behaviors;

Design mentalistic agents, inspired by agency's strong definition;

Enable real-time decision-making, to cope with dynamic and unpredictable environments;



# System Architecture




The system's modularity separates environment and agents in the MAS, using agent-oriented programming to give robots the autonomy to achieve their goals.

## ➤ Environment Class

- **House:** Simulated layout with defined rooms, obstacles, and openings.
- **Survivors:** Individuals needing rescue, placed in the environment.
- **Ambulance:** Vehicle transporting survivors to safety;

## ➤ Agent Class

- Agents behave using the **Belief-Desire-Intention** (BDI) framework.
- Use **distributed aggregative optimization** algorithm to make optimal decisions aiming to achieve both local and global tasks;
- Navigate the environment while performing **obstacle avoidance**;



Agent(s) Model: Russel and Norvig, 2002


# Consensus Theory

Consensus Theory provides mathematical tools to reach **agreement** among distributed agents on a quantity.

- **Mathematical Model:** agents  $i$  in graph  $G$  iteratively update their states via local communication with out-neighbours  $j \in \mathcal{N}_i^{out}$  (agents receiving data from  $i$ ):

$$z_i^{k+1} = \sum_{j \in \mathcal{N}_i^{out}} a_{ij} z_j^k$$

- **Adjacency Matrix:** weights  $a_{ij} \geq 0$  represent communication strength between  $i$  and  $j$  in directed graph  $G$ . For a fixed topology, it forms an *LTI* system  $z^{k+1} = Az^k$ , where  $A \in \mathbb{R}^{N \times N}$  is the square adjacency matrix of the communication network;
- **Average Consensus:** agents reach *consensus* on the average of the initial estimates if  $A$  is doubly stochastic (i.e.,  $A\mathbf{1} = \mathbf{1}$  and  $\mathbf{1}^T A = \mathbf{1}^T$ ) and  $G$  is strongly connected and aperiodic. This ensures *simple stability* of the *LTI* system;



Graph Connectivity, 5 Agents Network

# Optimization Theory

In MAS, optimization is the key to determine the most efficient agent trajectories from available alternatives:

- **Gradient Descent:** iteratively updates decision variables  $z^k \forall k \in \{1, \dots, max\_iters\}$  in the steepest descent direction  $d^k = -\nabla \ell(z^k)$  with step-size  $\alpha^k > 0$ , minimizing the cost  $\ell \in \mathbb{R}^d \rightarrow \mathbb{R}$  to solve  $\min_{z \in \mathbb{R}^d} \ell(z)$ :

$$z^{k+1} = z^k - \alpha^k \nabla \ell(z^k)$$

- **Convexity:** any local minimum of  $\ell(z)$  is also a global minimum, i.e.,  $z_i^*$  is an optimal global solution if:

$$\ell(\theta z_1 + (1 - \theta)z_2) \leq \theta \ell(z_1) + (1 - \theta)\ell(z_2) \forall z_1, z_2 \in \mathbb{R}^d, \theta \in [0,1];$$

- **Convergence:** convergence of  $z_i^k$  to  $z_i^* \forall i$  is guaranteed when  $\ell$  is convex and  $\alpha^k$  is sufficiently small to ensure a stable decrease in the cost, i.e.,  $\ell(z^{k+1}) < \ell(z^k) \forall k$  to solutions  $z^*$  such that  $\nabla \ell(z^*) = 0$  and  $\nabla^2 \ell(z^*) \geq 0$ ;

# Gradient Tracking Algorithm

- **Distributed Optimization:** Through local updates, agents aim to track a global estimate of the gradient while optimizing local decision estimates  $z_i^k$  of  $z_i^*$  via minimization of aggregated cost function  $\ell(z) = \sum_{i=1}^N \ell_i(z_i)$ ;
- **Dynamic Average Consensus:** each agent maintains a local estimate  $s_i^k$  (tracker) of a time-varying signal  $\overline{r^k}$ , (e.g., gradient and formation's barycenter), through an averaging system with local innovation to track changes;
- **Gradient Tracking Algorithm:** combines consensus and optimization to track total gradient  $\frac{1}{N} \sum_{i=1}^N \nabla \ell_i(z_i^k)$ :

$$z_i^{k+1} = \sum_{j \in \mathcal{N}_i^{out}} a_{ij} z_j^k - \alpha^k s_i^k, \quad z_i^0 \in \mathbb{R}^d \quad \forall i \in \{1, \dots, N\}$$
$$s_i^{k+1} = \sum_{j \in \mathcal{N}_i^{out}} a_{ij} s_j^k + \left( \nabla \ell_i(z_i^{k+1}) - \nabla \ell_i(z_i^k) \right), \quad s_i^0 = \nabla \ell(z_i^0) \quad \forall i \in \{1, \dots, N\},$$

- **Convergence:** the algorithm converges if  $A$  is doubly stochastic, graph  $G$  is connected and each cost function  $\ell_i$  is strongly convex with a Lipschitz gradient, ensuring  $\lim_{k \rightarrow \infty} \|z_i^k - z_i^*\| = 0$  for all agents  $i$ ;

# Distributed Aggregative Optimization

Aggregative Optimization involves  $N$  **robots** optimizing positions  $z_i \in \mathbb{R}^2, i \in 1, \dots, N$  by minimizing local cost functions  $\ell_i(z_i, \sigma(z))$ , which rely on global time-varying information estimated through Gradient Tracking algorithm:

- Agents  $i \in [1, N]$  independently optimize **local decision variables**  $z_i^k$ , i.e.,  $N$  different solutions  $z_i^*$  are computed;
- The problem involves **two global terms**: the cumulative gradient  $\sum_{j=1}^N \nabla_2 \ell_j(z_j^k, \sigma(z^k))$  and the barycenter position  $\sigma(z)$ ;
- Each agent  $i$  is aware only of  $\phi_i$  and  $\ell_i$ , and maintains **estimate**  $z_i^k$  of the optimal solution  $z_i^*$  and **trackers**  $s_i^k$  and  $v_i^k$  of global time-varying signals  $\sigma(z^k)$  and  $\sum_{j=1}^N \nabla_2 \ell_j(z_j^k, \sigma(z^k))$  to track;

$$\min_{z_1, \dots, z_N} \sum_{i=1}^N \ell_i(z_i, \sigma(z)), \quad \text{with} \quad \sigma(z) = \frac{1}{N} \sum_{i=1}^N \phi_i(z_i)$$

$$z_i^{k+1} = z_i^k - \alpha \left( \nabla_1 \ell_i(z_i^k, s_i^k) + \nabla \phi_i(z_i^k) v_i^k \right)$$

$$s_i^{k+1} = \sum_{j \in \mathcal{N}_i} a_{ij} s_j^k + \phi_i(z_i^{k+1}) - \phi_i(z_i^k)$$

$$v_i^{k+1} = \sum_{j \in \mathcal{N}_i} a_{ij} v_j^k + \nabla_2 \ell_i(z_i^{k+1}, s_i^{k+1}) - \nabla_2 \ell_i(z_i^k, s_i^k)$$

$$z_i^0 \in \mathbb{R}^{n_i}, s_i^0 = \phi_i(z_i^0), \text{ and } v_i^0 = \nabla_2 \ell_i(z_i^0, s_i^0)$$




# Obstacle Avoidance

- Integrated into agents' costs to navigate the environment while safely **avoiding collisions** and maintaining **tight formation**;
- The system assigns **potential functions** to obstacle points, generating **repulsive forces** with gain  $K$  as agents approach;
- They represent **artefacts** for the robotic agents in the MAS;
- The repulsive force, derived from the potential's gradient, directs agents along the **steepest descent** path of the function  $U_{rep}$ ;
- **Local minima** may arise in complex cluttered environments, potentially trapping agents in sub-optimal positions;

$$U_{\text{rep},i}(q) = \begin{cases} \frac{1}{2}K_{r_i} \left( \frac{1}{d_i(q)} - \frac{1}{q^*} \right)^2 & \text{if } d_i(q) < q^* \\ 0 & \text{otherwise} \end{cases}$$

$$\nabla U_{\text{rep},i}(q) = \begin{cases} K_{r_i} \left( \frac{1}{q^*} - \frac{1}{d_i(q)} \right) \frac{\nabla d_i(q)}{d_i(q)^2} & \text{if } d_i(q) \leq q^* \\ 0 & \text{if } d_i(q) > q^* \end{cases}$$



### Example of Room Walls Avoidance

# Cost Function


The algorithm's effectiveness hinges on the cost function convex structure, dictating **agents' behaviors**:

$$\ell_i(z_i, \sigma(z)) = \gamma_{rlt} \|z_i - r_i\|^2 + \gamma_{bar} \|\sigma(z) - b\|^2 + \gamma_{agg} \|z_i - \sigma(z)\|^2 + obst\_pot$$

- **Local Target Attraction:**  $\gamma_{rlt} \|z_i - r_i\|^2$ , based on the distance of each robot  $i$  to its local target  $r_i \in \mathbb{R}^2$ ;
- **Barycenter Goal:**  $\gamma_{bar} \|\sigma(z) - b\|^2$ , based on the distance between the barycenter  $\sigma(z)$  and global target  $b \in \mathbb{R}^2$ ;
- **Barycenter Attraction:**  $\gamma_{agg} \|z_i - \sigma(z)\|^2$ , based on the distance between the robot  $i$  and the barycenter  $\sigma(z)$ ;
- **Potential Function:**  $obst\_pot$ , based on the distance between the robot  $i$  and each obstacle point;


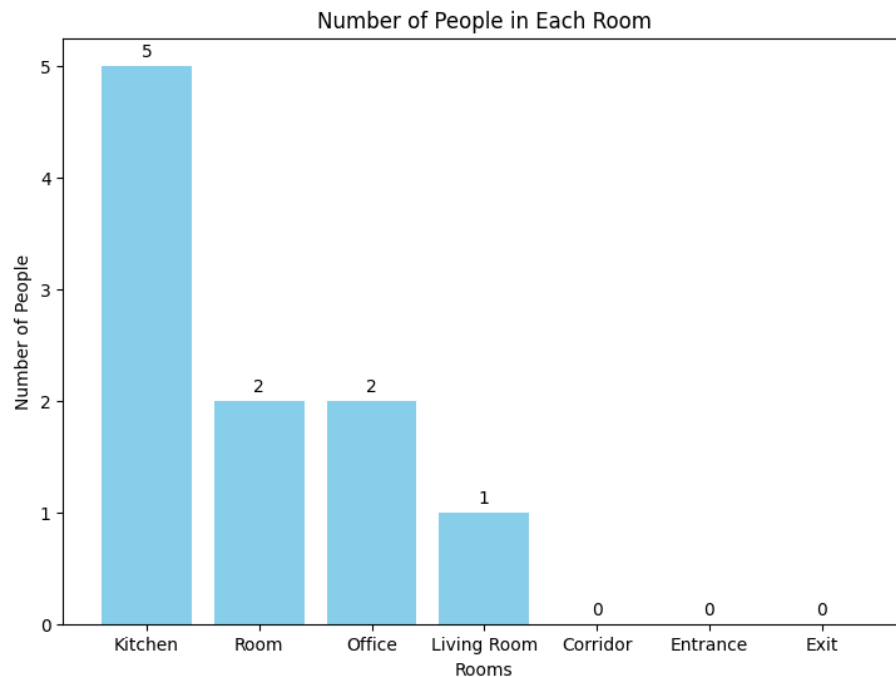
# Environment

- **Ambulance and Survivors:** FSMs (*artefacts*) used to simulate people in rooms and the ambulance;
- **Obstacles and Openings:** walls, obstacle points and openings define the navigable space for the agents;
- **Midpoints:** used for optimized navigation in plans, guiding agents efficiently between rooms;
- **Pragmatical Actions:** situated agents may change the state of the environment with pragmatical actions;
- **Visualization:** the *plot\_house()* function provides a color-coded graphical representation of the environment;



# Rooms Initialization

At the beginning of the simulation, survivors are randomly placed in rooms according to a **uniform distribution**. The exact number of people for each room is determined by the probabilities, and it is not known by the agents.



# Survivors




## Finite State Machine

Survivors are managed by an FSM with states: *Steady*, *Escorted*, *Saved*, *Transported* and *Healed*;

Randomly initialized in rooms according to house's heatmap;

*Saved* survivors use P-controllers to reach the *Idle* ambulance;



## Groups Creation

Agents gather *Steady* survivors using the *create\_escort\_group()* function;

*max\_survivors\_escort* caps the number of people escorted at once;

If survivors are left due to group limits, agents prioritize returning to their room over exploring new areas;



## Rescue

Survivors transition to *Escorted* when grouped and moved to *Exit* with agents, where they are *Saved*;

*Saved* survivors are delivered to the hospital upon ambulance arrival, entering *Transported* state;

The mission continues until all survivors are *Healed* in hospital;



# Ambulance



## Finite State Machine

Operate in three states: *Idle*,  
*Departing* and *Returning*;

Ambulance movements are  
synchronized with agents and  
survivors' activities;



## Arrival and Departure

*Departing* once a group of escorted  
survivors boards the ambulance;

After a while, another ambulance is  
dispatched to the exit in *Returning*  
state to collect the next group;




## Hospital

Transport survivors to the hospital,  
marking them *Healed* upon arrival;

If all survivors have been healed,  
then the ambulance must stay *Idle*,  
to prevent unnecessary dispatches;

# Beliefs, Desires, Intentions



## Beliefs

Represent agent's knowledge about the MAS environment;

Initialized from a YAML file;

Include room locations, obstacle points and survivor statuses;

Dynamically updated through *perceive\_environment()* and *brf()*;



## Desires

Define the agent's objectives based on the current context, such as escorting survivors, visiting rooms or reach exit;

Function *generate\_options()* generates feasible desires on updated beliefs;

Multiple, possibly conflicting, desires can co-exist inside agents' minds;



## Intentions

Represent the plans and actions chosen to fulfill agents' desires;

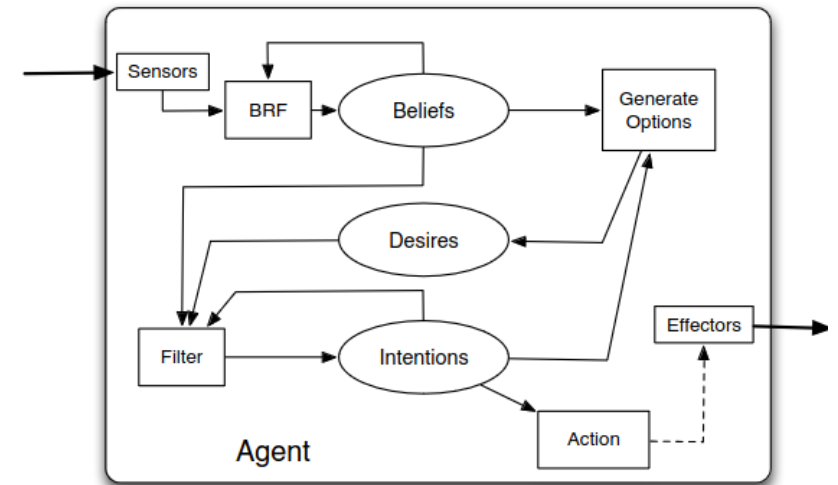
Function *filter\_options()* selects and prioritizes one single desire;

Implemented by *create\_plan()* through waypoints, obstacles and escort groups' generation;

# BDI Control Loop

The control loop of BDI agents represents their “**minds**”, continuously updating beliefs, generating desires, filtering options into intentions and executing plans to achieve their goals. Such a model follows agency’s strong definition.

```
while True:
    perceptions = perceive_environment()
    beliefs = belief_revision_function(perceptions, beliefs)
    desires = generate_options(beliefs)
    intentions = filter_options(beliefs, desires)
    plan, obstacles = create_plan(beliefs, intentions)
    execute_plan(plan, obstacles, beliefs, desires, intentions)
    if check_termination(): break
```

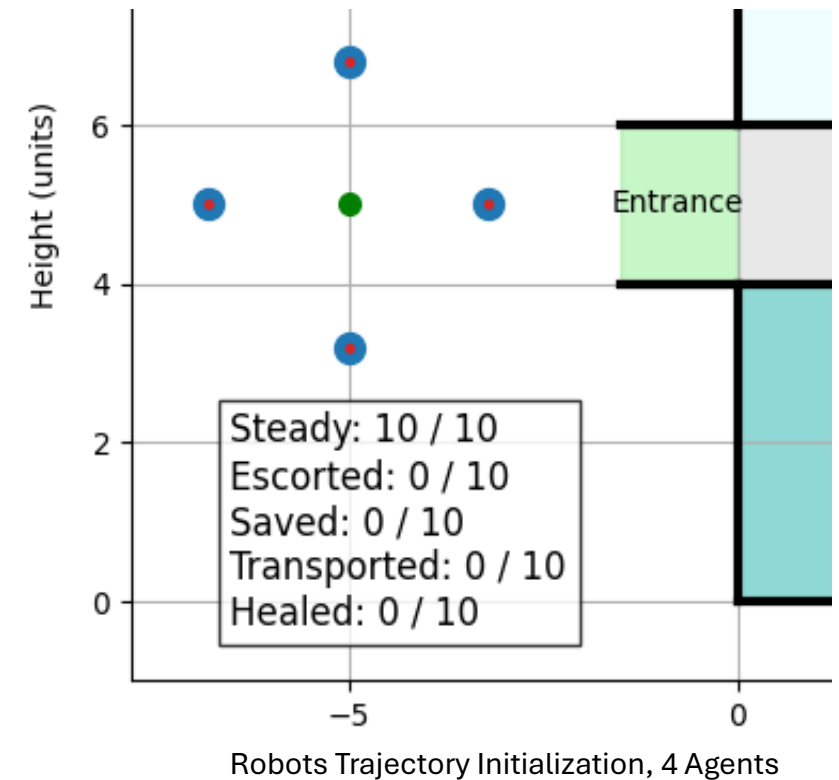


Basic Architecture of a BDI Agent, Wooldridge 2009

# Trajectories Initialization

The *trajectory\_initialization()* method sets up the trajectories for the MAS' components across iterations, space and dimensions.

- **Trajectory Initialization:** prepares trajectories for formation targets ( $R, b$ ), agent positions ( $Z$ ), barycenter and gradient trackers ( $S, V$ );
- **Circular Formation:** utilize *generate\_circular\_positions()* to set the robots' initial formation radially around entrance point with  $N$  agents;
- **Targets Definition:** similarly to circular formation, set local target coordinates  $R$  radially scattered around global coordinates  $b$ ;
- **Cost and Gradient Monitoring:** initialize arrays to track the cost function  $F$  and gradients ( $grad_z, grad_s$ ), essential for movements;



# Techniques and Procedures



## Entering the House

Agents initially move to the local targets placed in *Entrance* room;  
Visit rooms with highest likelihood of finding survivors, continuously updating beliefs and desires;  
A plan is created to fulfill the intention, using rooms' waypoints, obstacles and formation's gains;



## Escorting Survivors

Agents form a feasible escort group once they detect survivors in rooms;  
Safely navigate to the exit adjusting their movements to avoid obstacles while maintaining a tight formation;  
Perform communication and pragmatical actions with neighbors and the environment in the MAS;



## Exiting the House

Once all known survivors have left the rooms, set intention to *Exit*;  
Agents may enter *Standby* mode at the exit if all survivors are *Saved*, but not yet fully *Healed*;  
Agents in the MAS are driven by the *goal* of ensuring all survivors are safely *Escorted* and *Healed*;

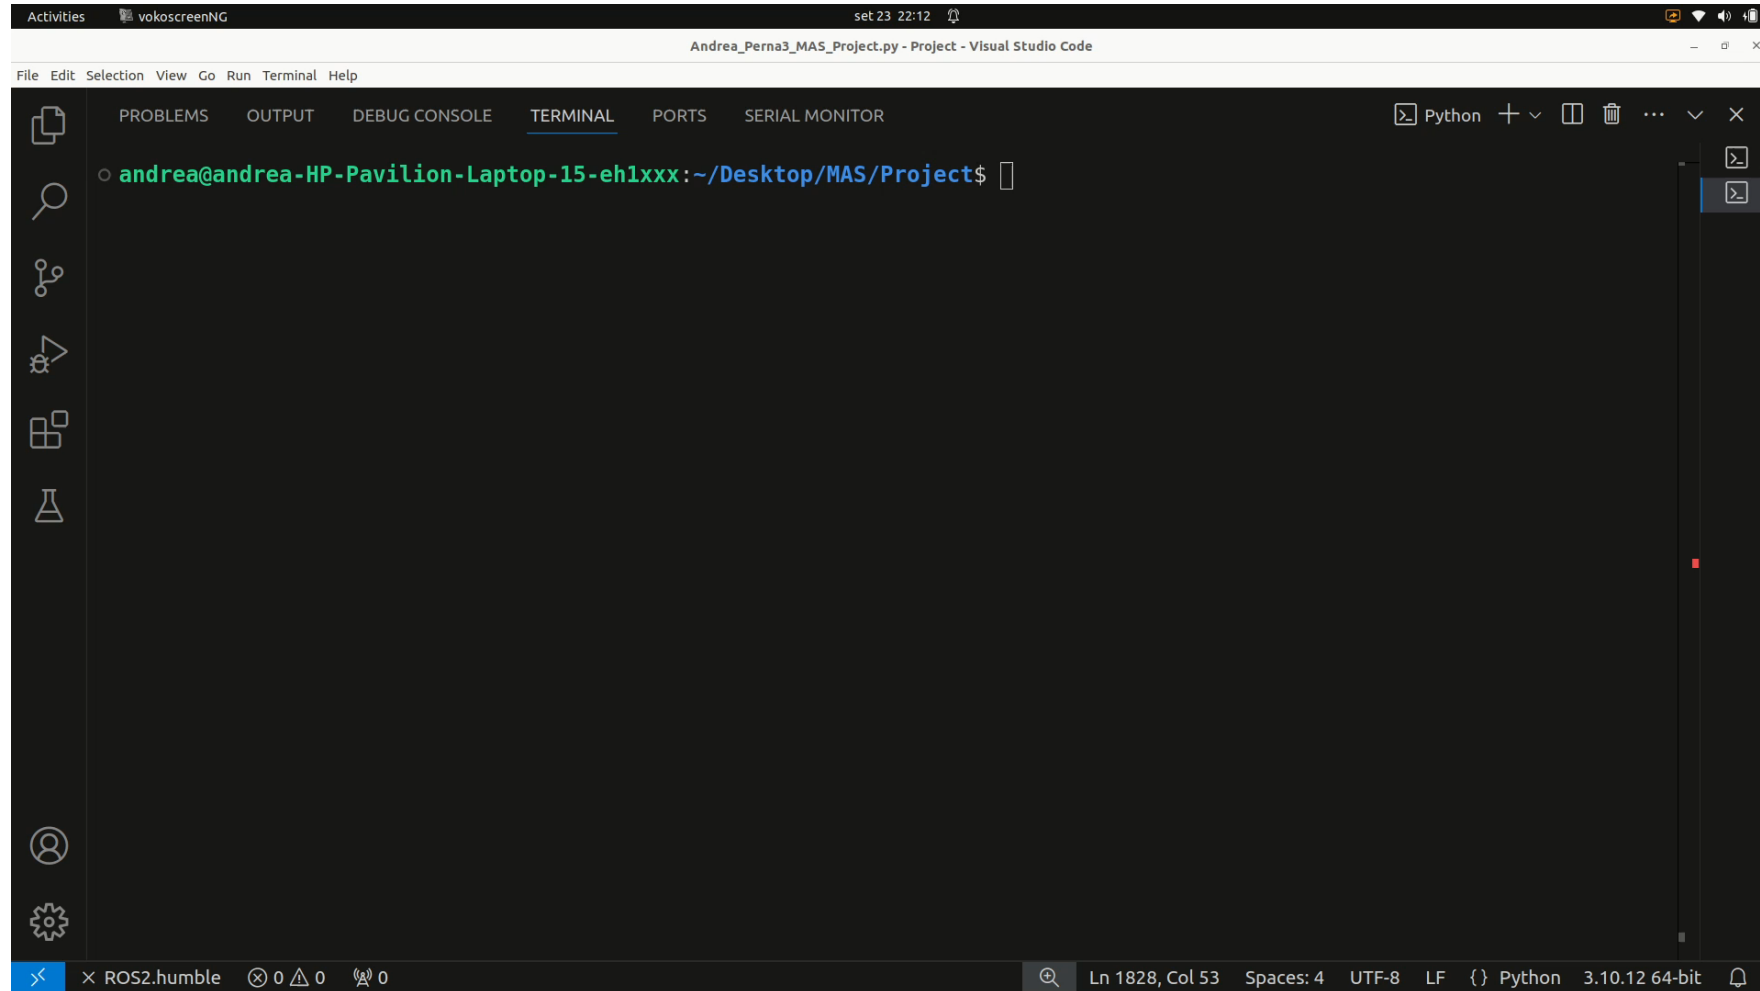


# Visualization


The *SAR\_animation()* function animates the trajectories of robots, survivors and the ambulance across iterations.

- **Plot Limits:** based on *view\_type* parameter in YAML file, they can be either static or dynamic;
- **House Layout:** house is shown in background with walls, corridors, obstacles in a color-coded fashion;
- **Robots Trajectories:** robots' initial positions and trajectories are shown to illustrate movement strategies;
- **Survivors Trajectories:** people's positions are dynamically updated, showing progress from detection to rescue;
- **Ambulance Trajectory:** the ambulance, with details like blinking lights, visually reflects its current state;
- **Dynamic Annotations:** real-time updates on key metrics, like iterations left and survivor status;

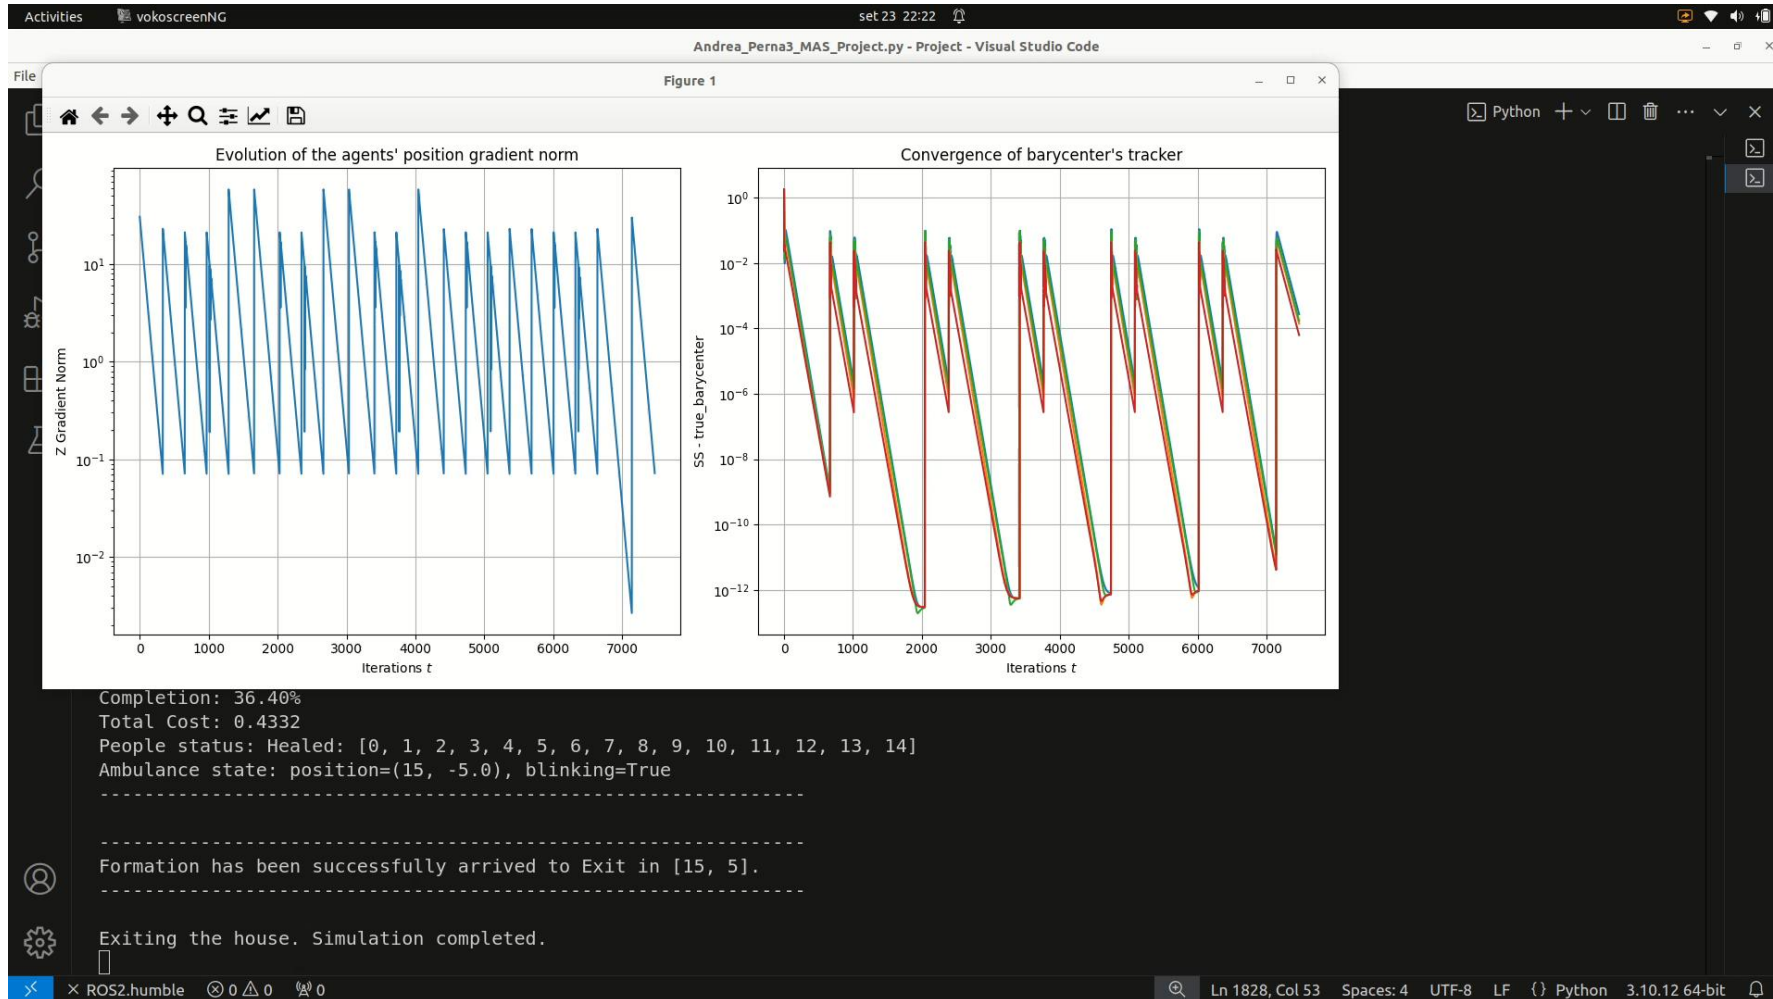
# SAR Animation – Computations



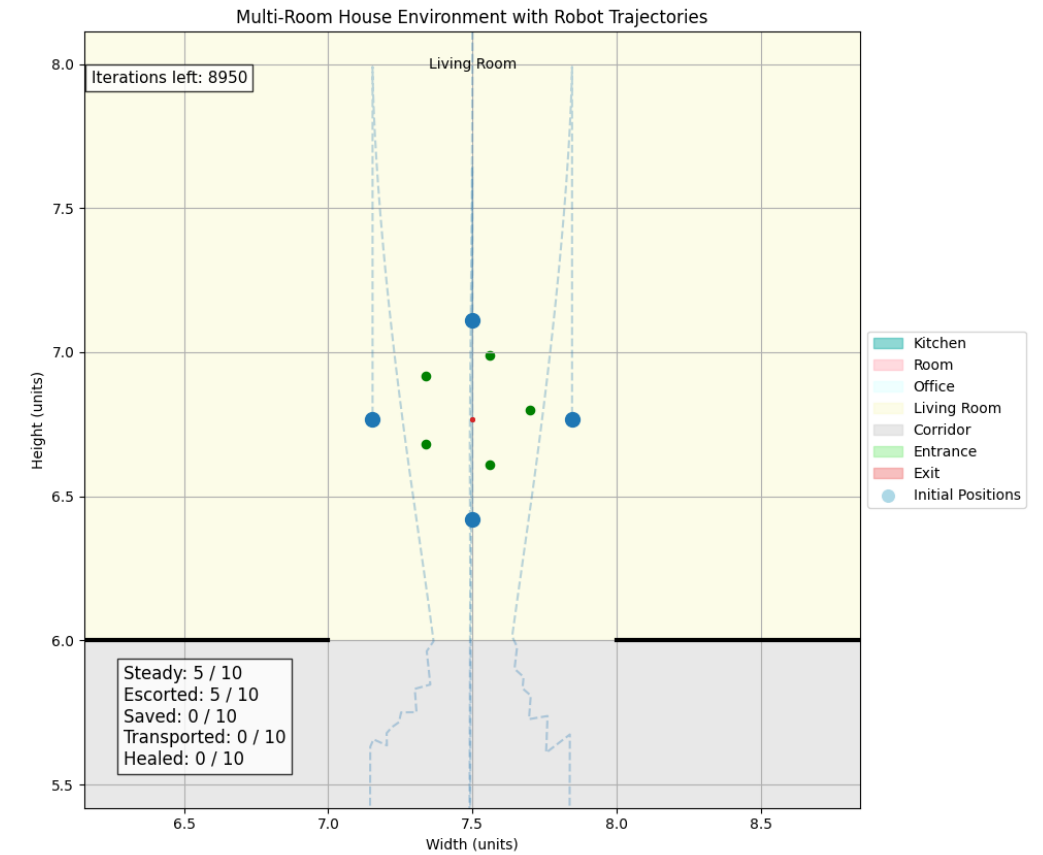

# SAR Animation – Static View



# SAR Animation – Static View

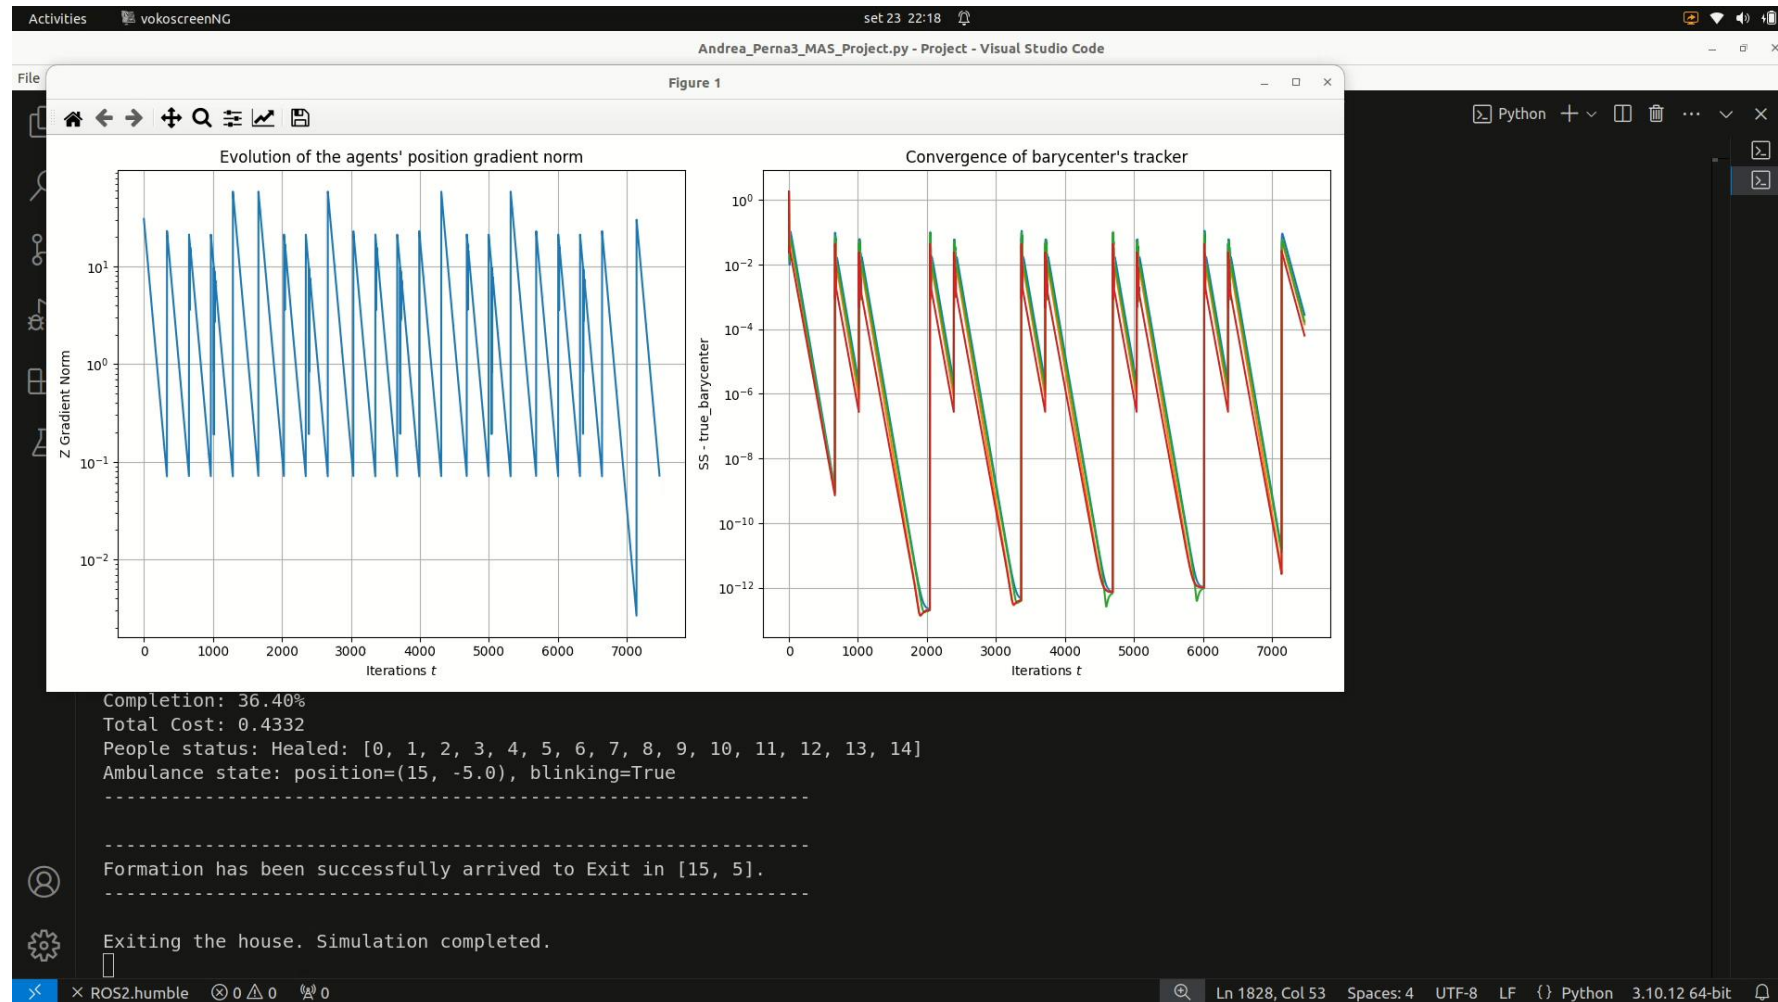


# SAR Animation – Dynamic View







# SAR Animation – Dynamic View



# Costs and Gradients



# Summary


To summarize, the **main characteristics** of the developed Multi-Agent System project are:

- Modularity and Extensibility;
- Distributed Aggregative Optimization;
- BDI Mentalistic Agents;
- Obstacle Avoidance via Potential Functions;
- Real-time Consensus Coordination;
- Visualization of SAR Simulations;



# Future Enhancements

- Implementation in **ROS2** distributed framework;
- Simultaneous Localization and Mapping (**SLAM**)
- Multi-Agent Reinforcement Learning (**MARL**);
- Tuple-based Coordination (**TuCSon**);
- **Battery** management and **energy** constraints;
- Decentralized **task allocation** to handle resources;
- Keep track of **health statuses** of survivors;
- **Computer Vision** features to enhance robots' perception.



# Thanks for your attention!



*Keep The Gradient*

- Bruno Siciliano -